

AmanAgrahari1314 / Speech_Emotion_Recognition

Code Issues Pull requests Actions Projects Wiki Security Insights Settings

Files main Go to file Speech Emotion Recognition - S... dataset

Speech Emotion Recognition / Speech Emotion Recognition - Sound Classification.ipynb

AmanAgrahari1314 Add files via upload 412f0f1 · 3 minutes ago History

Preview Code Blame 1390 lines (1390 loc) · 1.74 MB Code 55% faster with GitHub Copilot Raw

Import Modules

```
In [30]: import pandas as pd
import numpy as np
import os
import seaborn as sns
import matplotlib.pyplot as plt
import librosa
import librosa.display
from IPython.display import Audio
import warnings
warnings.filterwarnings('ignore')
```

Load the Dataset

```
In [31]: paths = []
labels = []
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        paths.append(os.path.join(dirname, filename))
        label = filename.split('.')[1]
        label = label.lower()
        labels.append(label)
if len(paths) == 2800:
    break
print('Dataset is loaded')

Dataset is loaded
```

```
In [32]: len(paths)
```

```
Out[32]: 2800
```

```
In [33]: paths[:5]
```

```
Out[33]: ['/kaggle/input/toronto-emotional-speech-set-tess/TESS Toronto emotional speech set data/YAF_fear/YAF_home_fear.wav',
'/kaggle/input/toronto-emotional-speech-set-tess/TESS Toronto emotional speech set data/YAF_fear/YAF_youth_fear.wav',
'/kaggle/input/toronto-emotional-speech-set-tess/TESS Toronto emotional speech set data/YAF_fear/YAF_near_fear.wav',
'/kaggle/input/toronto-emotional-speech-set-tess/TESS Toronto emotional speech set data/YAF_fear/YAF_search_fear.wav',
'/kaggle/input/toronto-emotional-speech-set-tess/TESS Toronto emotional speech set data/YAF_fear/YAF_pick_fear.wav']
```

```
In [34]: labels[:5]
```

```
Out[34]: ['fear', 'fear', 'fear', 'fear', 'fear']
```

```
In [35]: ## Create a dataframe
df = pd.DataFrame()
df['speech'] = paths
df['label'] = labels
df.head()
```

```
Out[35]: speech label
0 /kaggle/input/toronto-emotional-speech-set-tess/... fear
1 /kaggle/input/toronto-emotional-speech-set-tess/... fear
2 /kaggle/input/toronto-emotional-speech-set-tess/... fear
3 /kaggle/input/toronto-emotional-speech-set-tes... fear
4 /kaggle/input/toronto-emotional-speech-set-tes...
```

```
In [36]: df['label'].value_counts()
```

```
Out[36]: fear      400
angry     400
disgust   400
neutral   400
sad       400
ps        400
happy     400
Name: label, dtype: int64
```

Exploratory Data Analysis

```
In [37]: sns.countplot(data=df, x='label')
```

```
Out[37]: <AxesSubplot:xlabel='label', ylabel='count'>
```

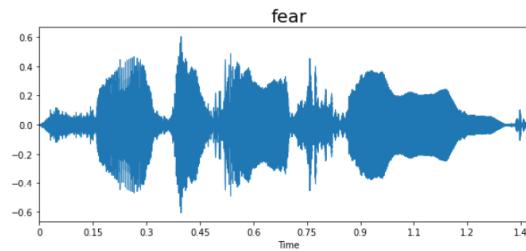
```
In [38]: def waveplot(data, sr, emotion):
    plt.figure(figsize=(10,4))
    plt.title(emotion, size=20)
    librosa.display.waveplot(data, sr=sr)
    plt.show()

def spectrogram(data, sr, emotion):
    # Your implementation here
```

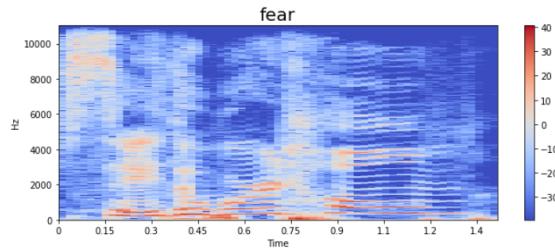
```
^ = librosa.core.spectrum.waveform
xdb = librosa.amplitude_to_db(abs(x))
plt.figure(figsize=(11,4))
plt.title(emotion, size=20)
librosa.display.specshow(xdb, sr=sr, x_axis='time', y_axis='hz')
plt.colorbar()
```

```
In [39]:
```

```
emotion = 'fear'
path = np.array(df['speech'][[df['label']==emotion]])[0]
data, sampling_rate = librosa.load(path)
waveplot(data, sampling_rate, emotion)
spectrogram(data, sampling_rate, emotion)
Audio(path)
```

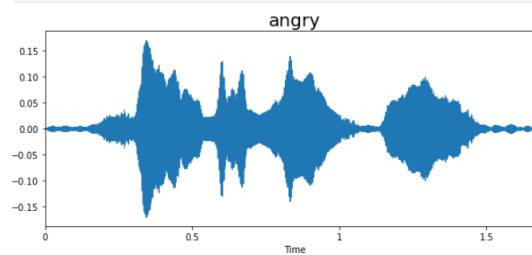


```
Out[39]:
```

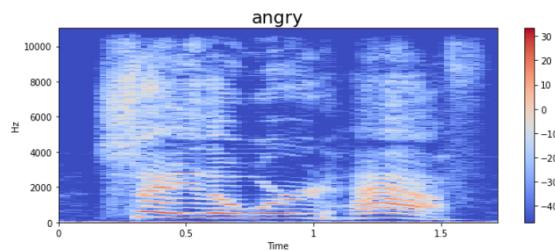


```
In [40]:
```

```
emotion = 'angry'
path = np.array(df['speech'][[df['label']==emotion]])[1]
data, sampling_rate = librosa.load(path)
waveplot(data, sampling_rate, emotion)
spectrogram(data, sampling_rate, emotion)
Audio(path)
```

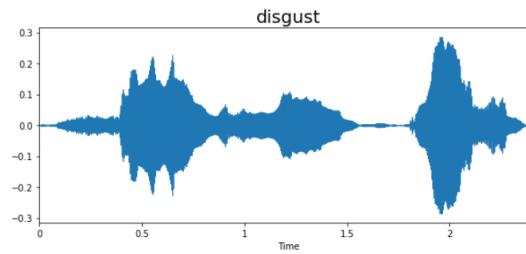


```
Out[40]:
```



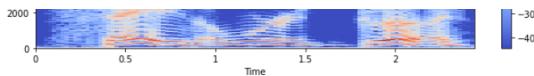
```
In [41]:
```

```
emotion = 'disgust'
path = np.array(df['speech'][[df['label']==emotion]])[0]
data, sampling_rate = librosa.load(path)
waveplot(data, sampling_rate, emotion)
spectrogram(data, sampling_rate, emotion)
Audio(path)
```

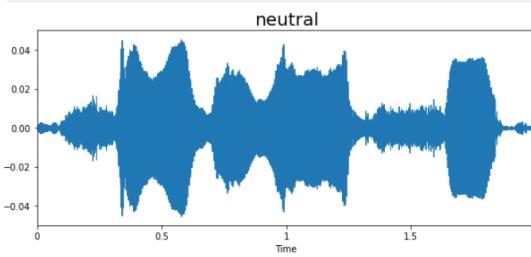


```
Out[41]:
```

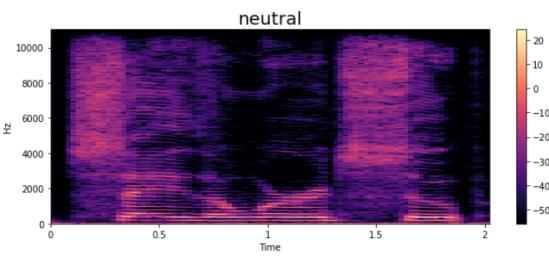




```
In [42]: emotion = 'neutral'  
path = np.array(df['speech'][df['label']==emotion])[0]  
data, sampling_rate = librosa.load(path)  
waveplot(data, sampling_rate, emotion)  
spectrogram(data, sampling_rate, emotion)  
Audio(path)
```

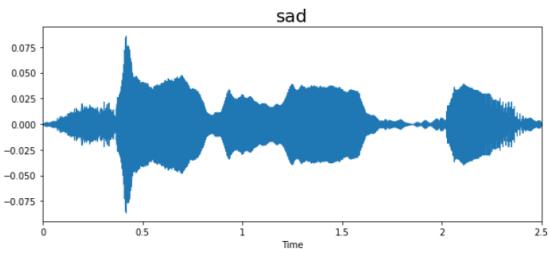


Out[42]:

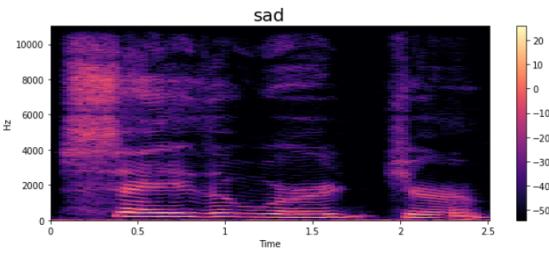


```
In [43]:
```

```
emotion = 'sad'  
path = np.array(df['speech'][df['label']==emotion])[0]  
data, sampling_rate = librosa.load(path)  
waveplot(data, sampling_rate, emotion)  
spectrogram(data, sampling_rate, emotion)  
Audio(path)
```

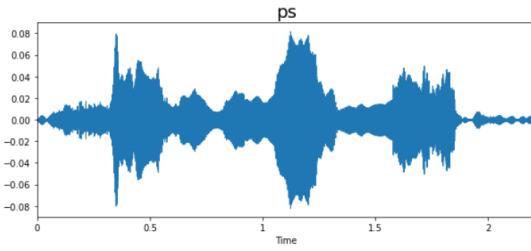


Out[43]:

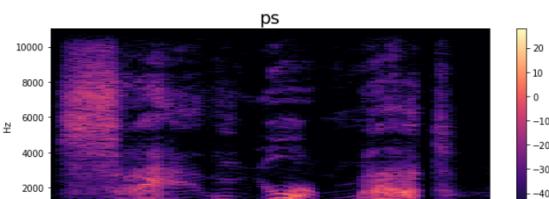


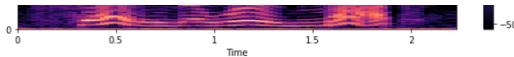
```
In [44]:
```

```
emotion = 'ps'  
path = np.array(df['speech'][df['label']==emotion])[0]  
data, sampling_rate = librosa.load(path)  
waveplot(data, sampling_rate, emotion)  
spectrogram(data, sampling_rate, emotion)  
Audio(path)
```

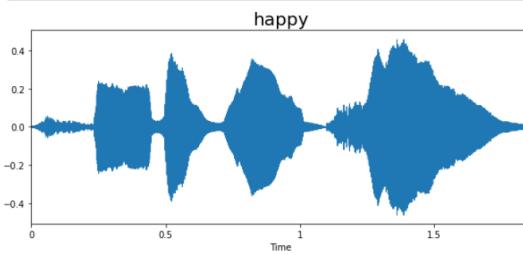


Out[44]:

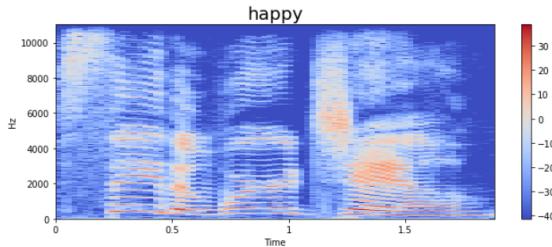




```
In [45]: emotion = 'happy'
path = np.array(df['speech'][df['label']==emotion])[0]
data, sampling_rate = librosa.load(path)
waveplot(data, sampling_rate, emotion)
spectrogram(data, sampling_rate, emotion)
Audio(path)
```



Out[45]:



Feature Extraction

```
In [46]: def extract_mfcc(filename):
    y, sr = librosa.load(filename, duration=3, offset=0.5)
    mfcc = np.mean(librosa.feature.mfcc(y=y, sr=sr, n_mfcc=40).T, axis=0)
    return mfcc
```

```
In [47]: extract_mfcc(df['speech'][0])
```

```
Out[47]: array([-285.2542 ,  86.24267 , -2.7735834 ,  22.61731 ,
 -15.214631 , 11.602871 , 11.931779 , -2.5318177 ,
 0.65986294, 11.62756 , -17.814924 , -7.5654893 ,
 6.2167853 , -5.7255652 , -9.565386 , 3.899267 ,
 -13.657834 , 14.420068 , 19.243341 , 23.024492 ,
 32.129776 , 16.585697 , -4.137755 , 1.2746525 ,
 -11.517016 , 7.0145273 , -2.8494127 , -7.415011 ,
 -11.150621 , -2.119054 , -5.4515266 , 4.473824 ,
 -11.377713 , -8.931878 , -3.8482994 , 4.950999 ,
 -1.7254968 , 2.659218 , 11.390564 , 11.3327265 ],
 dtype=float32)
```

```
In [48]: X_mfcc = df['speech'].apply(lambda x: extract_mfcc(x))
```

```
In [49]: X_mfcc
```

```
Out[49]: 0   [-285.2542 ,  86.24267 , -2.7735834 ,  22.61731 , -1...
 1   [-348.23337, 35.60242 , -4.365128, 15.514869, 6...
 2   [-339.50309, 54.41241 , -14.795754, 21.566118, ...
 3   [-306.92944, 21.973307 , -5.1588626, 7.6269317, ...
 4   [-344.88586, 47.05694 , -24.83122 , 20.24460 , 1...
 ...
 2795  [-374.1317, 61.859463 , -0.41998756, 9.31088, ...
 2796  [-314.12222, 40.262197 , -6.7989045, -3.2963052...
 2797  [-357.65854, 78.49201 , -15.684815, 3.644915, ...
 2798  [-352.78336, 102.219765 , -14.560364, -11.48181...
 2799  [-389.80002, 54.120773 , -0.8988281, -0.6595729, ...
Name: speech, Length: 2800, dtype: object
```

```
In [50]: X = [x for x in X_mfcc]
X = np.array(X)
X.shape
```

Out[50]: (2800, 40)

```
In [51]: ## input split
X = np.expand_dims(X, -1)
X.shape
```

Out[51]: (2800, 40, 1)

```
In [52]: from sklearn.preprocessing import OneHotEncoder
enc = OneHotEncoder()
y = enc.fit_transform(df[['label']])
```

```
In [53]: y = y.toarray()
```

```
In [54]: y.shape
```

Out[54]: (2800, 7)

Create the LSTM Model

```
In [101]: from keras.models import Sequential
from keras.layers import Dense, LSTM, Dropout

model = Sequential([
    LSTM(256, return_sequences=False, input_shape=(40,1)),
    Dropout(0.2),
    Dense(128, activation='relu'),
    Dropout(0.2),
    Dense(64, activation='relu'),
    Dropout(0.2),
```

```

        Dense(7, activation='softmax')
    ])

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.summary()

Model: "sequential_24"
-----  

Layer (type)      Output Shape     Param #
-----  

lstm_30 (LSTM)   (None, 256)     264192  

dropout_62 (Dropout) (None, 256)     0  

dense_72 (Dense) (None, 128)     32896  

dropout_63 (Dropout) (None, 128)     0  

dense_73 (Dense) (None, 64)      8256  

dropout_64 (Dropout) (None, 64)      0  

dense_74 (Dense) (None, 7)       455  

-----  

Total params: 305,799  

Trainable params: 305,799  

Non-trainable params: 0
-----
```

```

In [102]: # Train the model
history = model.fit(X, y, validation_split=0.2, epochs=50, batch_size=64)

Epoch 1/50
35/35 [=====] - 2s 18ms/step - loss: 1.0892 - accuracy: 0.6201 - val_loss: 2.0684 - val_accuracy: 0.2946
Epoch 2/50
35/35 [=====] - 0s 7ms/step - loss: 0.3742 - accuracy: 0.8598 - val_loss: 2.4078 - val_accuracy: 0.2054
Epoch 3/50
35/35 [=====] - 0s 8ms/step - loss: 0.1671 - accuracy: 0.9487 - val_loss: 1.9055 - val_accuracy: 0.4446
Epoch 4/50
35/35 [=====] - 0s 8ms/step - loss: 0.1672 - accuracy: 0.9442 - val_loss: 2.7364 - val_accuracy: 0.3179
Epoch 5/50
35/35 [=====] - 0s 7ms/step - loss: 0.1107 - accuracy: 0.9683 - val_loss: 1.8414 - val_accuracy: 0.5607
Epoch 6/50
35/35 [=====] - 0s 8ms/step - loss: 0.1082 - accuracy: 0.9585 - val_loss: 2.7303 - val_accuracy: 0.4679
Epoch 7/50
35/35 [=====] - 0s 9ms/step - loss: 0.1275 - accuracy: 0.9656 - val_loss: 1.2245 - val_accuracy: 0.7232
Epoch 8/50
35/35 [=====] - 0s 7ms/step - loss: 0.0765 - accuracy: 0.9786 - val_loss: 2.9995 - val_accuracy: 0.3893
Epoch 9/50
35/35 [=====] - 0s 8ms/step - loss: 0.0577 - accuracy: 0.9795 - val_loss: 3.5068 - val_accuracy: 0.3679
Epoch 10/50
35/35 [=====] - 0s 8ms/step - loss: 0.0531 - accuracy: 0.9826 - val_loss: 2.3169 - val_accuracy: 0.4786
Epoch 11/50
35/35 [=====] - 0s 8ms/step - loss: 0.0357 - accuracy: 0.9888 - val_loss: 3.8930 - val_accuracy: 0.4321
Epoch 12/50
35/35 [=====] - 0s 7ms/step - loss: 0.0554 - accuracy: 0.9835 - val_loss: 1.5915 - val_accuracy: 0.6607
Epoch 13/50
35/35 [=====] - 0s 7ms/step - loss: 0.0684 - accuracy: 0.9790 - val_loss: 3.4805 - val_accuracy: 0.4589
Epoch 14/50
35/35 [=====] - 0s 7ms/step - loss: 0.0453 - accuracy: 0.9835 - val_loss: 2.4777 - val_accuracy: 0.4661
Epoch 15/50
35/35 [=====] - 0s 8ms/step - loss: 0.0874 - accuracy: 0.9737 - val_loss: 4.4744 - val_accuracy: 0.2446
Epoch 16/50
35/35 [=====] - 0s 8ms/step - loss: 0.0530 - accuracy: 0.9853 - val_loss: 2.9993 - val_accuracy: 0.5232
Epoch 17/50
35/35 [=====] - 0s 6ms/step - loss: 0.0421 - accuracy: 0.9839 - val_loss: 4.3298 - val_accuracy: 0.4714
Epoch 18/50
35/35 [=====] - 0s 8ms/step - loss: 0.0514 - accuracy: 0.9857 - val_loss: 2.2161 - val_accuracy: 0.5946
Epoch 19/50
35/35 [=====] - 0s 7ms/step - loss: 0.0310 - accuracy: 0.9897 - val_loss: 3.7546 - val_accuracy: 0.4071
Epoch 20/50
35/35 [=====] - 0s 7ms/step - loss: 0.0301 - accuracy: 0.9897 - val_loss: 2.7526 - val_accuracy: 0.5036
Epoch 21/50
35/35 [=====] - 0s 8ms/step - loss: 0.0342 - accuracy: 0.9875 - val_loss: 4.7068 - val_accuracy: 0.2839
Epoch 22/50
35/35 [=====] - 0s 8ms/step - loss: 0.0295 - accuracy: 0.9893 - val_loss: 3.4425 - val_accuracy: 0.4054
Epoch 23/50
35/35 [=====] - 0s 7ms/step - loss: 0.0213 - accuracy: 0.9933 - val_loss: 2.8260 - val_accuracy: 0.5607
Epoch 24/50
35/35 [=====] - 0s 7ms/step - loss: 0.0255 - accuracy: 0.9933 - val_loss: 4.4797 - val_accuracy: 0.4696
Epoch 25/50
35/35 [=====] - 0s 7ms/step - loss: 0.0236 - accuracy: 0.9915 - val_loss: 4.2527 - val_accuracy: 0.4143
Epoch 26/50
35/35 [=====] - 0s 7ms/step - loss: 0.0223 - accuracy: 0.9920 - val_loss: 3.5158 - val_accuracy: 0.4429
Epoch 27/50
35/35 [=====] - 0s 7ms/step - loss: 0.0245 - accuracy: 0.9929 - val_loss: 3.9560 - val_accuracy: 0.4661
Epoch 28/50
35/35 [=====] - 0s 7ms/step - loss: 0.0238 - accuracy: 0.9902 - val_loss: 4.4557 - val_accuracy: 0.3893
Epoch 29/50
35/35 [=====] - 0s 8ms/step - loss: 0.0702 - accuracy: 0.9772 - val_loss: 3.5628 - val_accuracy: 0.3839
Epoch 30/50
35/35 [=====] - 0s 8ms/step - loss: 0.0165 - accuracy: 0.9951 - val_loss: 3.8458 - val_accuracy: 0.4089
Epoch 31/50
35/35 [=====] - 0s 8ms/step - loss: 0.0080 - accuracy: 0.9978 - val_loss: 3.8509 - val_accuracy: 0.4339
Epoch 32/50
35/35 [=====] - 0s 7ms/step - loss: 0.0091 - accuracy: 0.9964 - val_loss: 3.8585 - val_accuracy: 0.4786
Epoch 33/50
35/35 [=====] - 0s 6ms/step - loss: 0.0146 - accuracy: 0.9960 - val_loss: 4.4246 - val_accuracy: 0.3554
Epoch 34/50
35/35 [=====] - 0s 7ms/step - loss: 0.0075 - accuracy: 0.9969 - val_loss: 4.4920 - val_accuracy: 0.3911
Epoch 35/50
35/35 [=====] - 0s 7ms/step - loss: 0.0072 - accuracy: 0.9982 - val_loss: 3.6941 - val_accuracy: 0.4232

```

```

tpeoch 30/50
35/35 [=====] - 0s 7ms/step - loss: 0.0244 - accuracy: 0.9933 - val_loss: 2.6108 - val_accuracy: 0.5
000
Epoch 37/50
35/35 [=====] - 0s 7ms/step - loss: 0.0115 - accuracy: 0.9969 - val_loss: 3.3635 - val_accuracy: 0.5
75
Epoch 38/50
35/35 [=====] - 0s 7ms/step - loss: 0.0519 - accuracy: 0.9853 - val_loss: 5.5903 - val_accuracy: 0.2
554
Epoch 39/50
35/35 [=====] - 0s 7ms/step - loss: 0.0369 - accuracy: 0.9906 - val_loss: 3.8724 - val_accuracy: 0.4
018
Epoch 40/50
35/35 [=====] - 0s 7ms/step - loss: 0.0122 - accuracy: 0.9964 - val_loss: 4.3779 - val_accuracy: 0.4
250
Epoch 41/50
35/35 [=====] - 0s 7ms/step - loss: 0.0124 - accuracy: 0.9973 - val_loss: 3.4232 - val_accuracy: 0.4
893
Epoch 42/50
35/35 [=====] - 0s 7ms/step - loss: 0.0095 - accuracy: 0.9969 - val_loss: 4.3362 - val_accuracy: 0.3
804
Epoch 43/50
35/35 [=====] - 0s 8ms/step - loss: 0.0089 - accuracy: 0.9978 - val_loss: 3.9718 - val_accuracy: 0.4
911
Epoch 44/50
35/35 [=====] - 0s 7ms/step - loss: 0.0058 - accuracy: 0.9987 - val_loss: 3.5679 - val_accuracy: 0.5
018
Epoch 45/50
35/35 [=====] - 0s 7ms/step - loss: 0.0074 - accuracy: 0.9982 - val_loss: 4.0037 - val_accuracy: 0.4
607
Epoch 46/50
35/35 [=====] - 0s 7ms/step - loss: 0.0030 - accuracy: 0.9991 - val_loss: 4.6531 - val_accuracy: 0.3
982
Epoch 47/50
35/35 [=====] - 0s 8ms/step - loss: 0.0076 - accuracy: 0.9982 - val_loss: 5.2379 - val_accuracy: 0.3
571
Epoch 48/50
35/35 [=====] - 0s 7ms/step - loss: 0.0082 - accuracy: 0.9973 - val_loss: 4.3685 - val_accuracy: 0.4
357
Epoch 49/50
35/35 [=====] - 0s 7ms/step - loss: 0.0155 - accuracy: 0.9964 - val_loss: 4.8508 - val_accuracy: 0.3
804
Epoch 50/50
35/35 [=====] - 0s 6ms/step - loss: 0.0079 - accuracy: 0.9982 - val_loss: 5.0355 - val_accuracy: 0.3
750

```

```

In [ ]: # best val accuracy: 72.32
# use checkpoint to save the best val accuracy model
# adjust learning rate for slow convergence

```

Plot the results

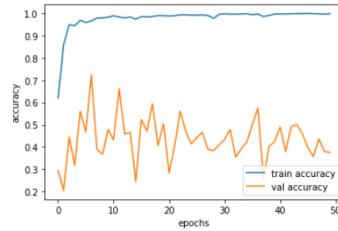
In [103...]

```

epochs = list(range(50))
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.plot(epochs, acc, label='train accuracy')
plt.plot(epochs, val_acc, label='val accuracy')
plt.xlabel('epochs')
plt.ylabel('accuracy')
plt.legend()
plt.show()

```



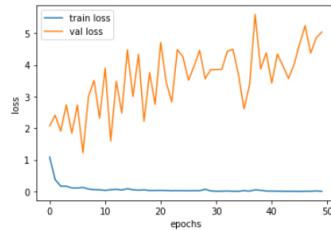
In [104...]

```

loss = history.history['loss']
val_loss = history.history['val_loss']

plt.plot(epochs, loss, label='train loss')
plt.plot(epochs, val_loss, label='val loss')
plt.xlabel('epochs')
plt.ylabel('loss')
plt.legend()
plt.show()

```



In []:

In []: